

# ETC3555 2018 - Lab 9 solutions

## Recommender Systems

*Cameron Roach and Souhaib Ben Taieb*

*11 October, 2018*

In this assignment, you will implement the collaborative filtering learning algorithm and apply it to a dataset of movie ratings. This dataset consists of ratings on a scale of 1 to 5. The dataset has  $n_u = 943$  users, and  $n_m = 1682$  movies.

```
load("movies_ratings.Rda")
list2env(data,.GlobalEnv) # "Y" and "R"

## <environment: R_GlobalEnv>

rm(data)

# Y is a 1682x943 matrix, containing ratings (1-5) of 1682 movies on
# 943 users
#
# R is a 1682x943 matrix, where R[i,j] <- 1 if and only if user j gave a
# rating to movie i
```

The following code gives all movies titles.

```
source("loadMovieList.R")
movieList <- loadMovieList()
```

(1 mark) Compute the average ratings for “Toy Story (1995)” and “Alaska (1996)”.

```
# ----- YOUR CODE HERE -----

library(tidyverse)

ifelse(R, Y, NA) %>%
  t() %>%
  magrittr::set_colnames(movieList) %>%
  as_data_frame() %>%
  select("Toy Story (1995)", "Alaska (1996)") %>%
  summarise_all(funs(mean(., na.rm=TRUE)))

## # A tibble: 1 x 2
##   `Toy Story (1995)` `Alaska (1996)`
##           <dbl>           <dbl>
## 1           3.88           2.69

# -----
```

(1 mark) Which user has rated more movies? Plot a histogram of her/his ratings?

```
# ----- YOUR CODE HERE -----

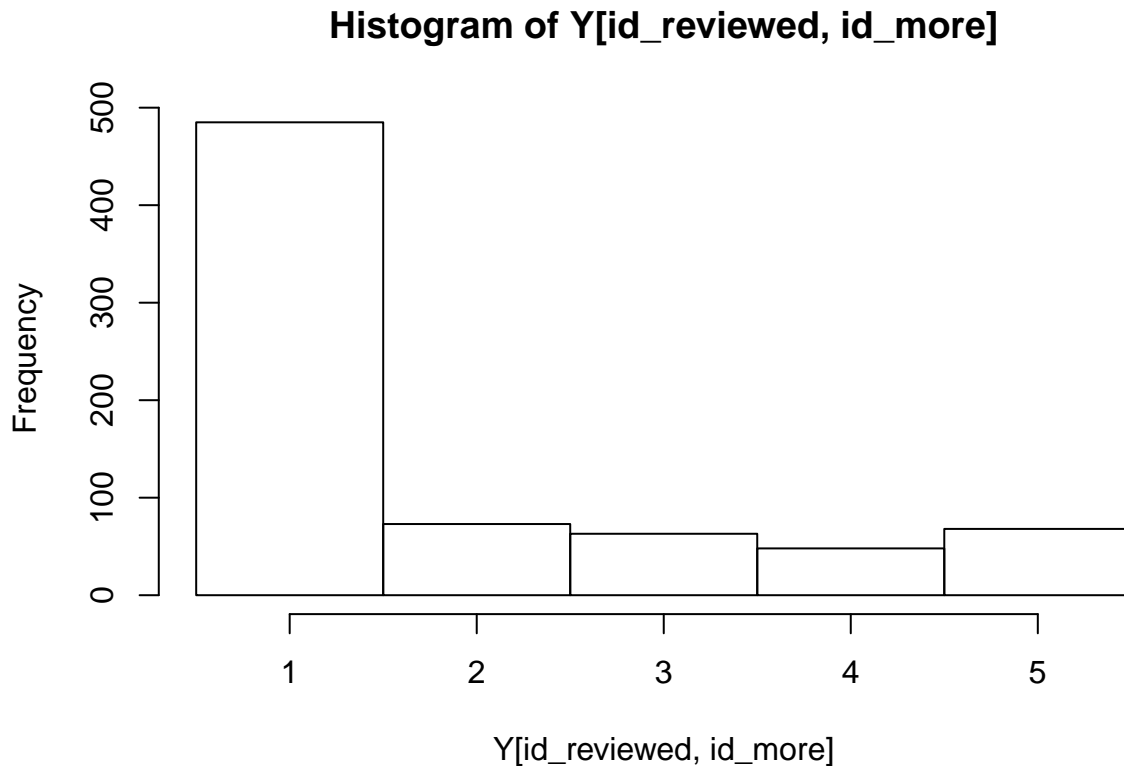
id_more <- which.max(apply(R, 2, sum))
id_reviewed <- which(R[, id_more] == 1)
print(id_more)

## [1] 405
```

```
print(length(Y[id_reviewed, id_more]))
```

```
## [1] 737
```

```
hist(Y[id_reviewed, id_more], .5+0:5)
```



```
# -----
```

You will now implement the cost function for collaborative filtering. Specifically, you should complete the following code to return J.

```
cofiCostFunc <- function(Y, R, num_users, num_movies,
                          num_features, lambda = 0) {
  #COFICOSTFUNC Collaborative filtering cost function
  # J <- COFICOSTFUNC(Y, R, num_users, num_movies, ...
  # num_features, lambda)(params) returns the cost for the
  # collaborative filtering problem.
  #
  # Notes: X - num_movies x num_features matrix of movie features
  #         Theta - num_users x num_features matrix of user features
  #         Y - num_movies x num_users matrix of user ratings of movies
  #         R - num_movies x num_users matrix, where R(i, j) <- 1 if the
  #            i-th movie was rated by the j-th user

  function(params) {
    # Unfold the U and W matrices from params
    X <-
      matrix(params[1:(num_movies * num_features)], num_movies, num_features)
    Theta <-
      matrix(params[(num_movies * num_features + 1):length(params)], num_users, num_features)
```

```

# ----- YOUR CODE HERE -----
J <- 0
J <- (1 / 2) * sum(((X %*% t(Theta)) * R - Y * R) ^ 2) +
  (lambda / 2 * sum(Theta ^ 2)) + (lambda / 2 * sum(X ^ 2))
# -----
J
}
}

```

To help you debug your cost function, run the following code.

```

# Load pre-trained weights (X, Theta, num.users, num.movies, num.features)
load("movieParams.Rda")
list2env(data, .GlobalEnv)

## <environment: R_GlobalEnv>

rm(data)

num_users <- as.numeric(num.users)
num_movies <- as.numeric(num.movies)
num_features <- as.numeric(num.features)

# Reduce the data set size so that this runs faster
num_users <- 4; num_movies <- 5; num_features <- 3
X <- X[1:num_movies, 1:num_features]
Theta <- Theta[1:num_users, 1:num_features]
Y <- Y[1:num_movies, 1:num_users]
R <- R[1:num_movies, 1:num_users]

# Evaluate cost function
J <- cofiCostFunc(Y, R, num_users, num_movies, num_features, 0)(c(X), c(Theta))

cat(sprintf('Cost at loaded parameters: %f (this value should be about 22.22)\n', J))

```

```
## Cost at loaded parameters: 22.224604 (this value should be about 22.22)
```

Once your cost function matches up with ours, you should now implement the collaborative filtering gradient function. Specifically, you should complete the following code to return the grad argument.

```

cofiGradFunc <- function(Y, R, num_users, num_movies, num_features, lambda = 0) {
  #cofiGradFunc returns the gradient for the
  # collaborative filtering problem.
  # Notes: X - num_movies x num_features matrix of movie features
  #         Theta - num_users x num_features matrix of user features
  #         Y - num_movies x num_users matrix of user ratings of movies
  #         R - num_movies x num_users matrix, where R[i, j] <- 1 if the
  #            i-th movie was rated by the j-th user
  #
  function(params) {
    # Unfold the U and W matrices from params
    X <-
      matrix(params[1:(num_movies * num_features)], num_movies, num_features)
    Theta <-
      matrix(params[(num_movies * num_features + 1):length(params)],
              num_users, num_features)
  }
}

```

```

# You should set the following variables correctly:
#       X_grad - num_movies x num_features matrix, containing the
#               partial derivatives w.r.t. to each element of X
#       Theta_grad - num_users x num_features matrix, containing the
#                   partial derivatives w.r.t. to each element of Theta
#
# ----- YOUR CODE HERE -----
X_grad <- matrix(0,dim(X)[1],dim(X)[2])
Theta_grad <- matrix(0, dim(Theta)[1], dim(Theta)[2])

X_grad <- (((X %*% t(Theta)) * R) %*% Theta - (Y * R) %*% Theta) + lambda * X
Theta_grad <- t((t(X) %*% ((X %*% t(Theta)) * R) - t(X) %*% (Y * R))) + lambda * Theta
# -----
grad <- c(c(X_grad),c(Theta_grad))
grad
}
}

```

You can check your function by running the following code.

```

source("computeNumericalGradient.R")
source("checkCostFunction.R")
checkCostFunction()

```

```

##          numgrad          grad
## [1,] -0.55499681 -0.55499681
## [2,]  0.69251864  0.69251864
## [3,] -2.42820624 -2.42820624
## [4,] -2.86158362 -2.86158362
## [5,]  4.36825009  4.36825009
## [6,]  2.13383447  2.13383447
## [7,]  2.06852264  2.06852264
## [8,]  0.03181708  0.03181708
## [9,]  2.43105156  2.43105156
## [10,] -2.80455561 -2.80455561
## [11,]  3.89199336  3.89199336
## [12,] -0.18570459 -0.18570459
## [13,] -0.56876254 -0.56876254
## [14,] -0.82111580 -0.82111580
## [15,]  3.48401697  3.48401697
## [16,] -1.36020376 -1.36020376
## [17,]  0.44965466  0.44965466
## [18,] -0.35655886 -0.35655886
## [19,] -1.84466173 -1.84466173
## [20,]  0.89889296  0.89889296
## [21,] -3.89654172 -3.89654172
## [22,]  2.34042388  2.34042388
## [23,] -0.73259714 -0.73259714
## [24,]  0.50044757  0.50044757
## [25,]  1.48144603  1.48144603
## [26,] -0.10139936 -0.10139936
## [27,] -0.84233864 -0.84233864

```

```
## The above two columns you get should be very similar.
##      (Left-Your Numerical Gradient, Right-Analytical Gradient)
##
## If your implementation is correct, then
##      the relative difference will be small (less than 1e-9).
##
##      Relative Difference: 1.08708e-12
```

(3 marks) Now, you should implement regularization for the cost function for collaborative filtering. Update the code of the *cofiCostFunc* function.

You can check your function using the following code.

```
lambda <- 1.5
J <- cofiCostFunc(Y, R, num_users, num_movies, num_features, lambda)(c(c(X),c(Theta)))

cat(sprintf('Cost at loaded parameters (lambda = 1.5): %f (this value should be about 31.34)\n', J))

## Cost at loaded parameters (lambda = 1.5): 31.344056 (this value should be about 31.34)
```

(3 marks) Once your cost matches up with ours, you should proceed to implement regularization for the gradient. Update the code of the *cofiGradFunc* function.

You can check the gradient computations with the following code.

```
checkCostFunction(lambda)

##          numgrad      grad
## [1,] -4.43276533 -4.43276533
## [2,] -1.86651963 -1.86651963
## [3,]  8.90525711  8.90525711
## [4,] -4.64843021 -4.64843021
## [5,] -2.71215493 -2.71215493
## [6,]  4.02522675  4.02522675
## [7,] -1.87156418 -1.87156418
## [8,]  4.70344519  4.70344519
## [9,] -2.46958225 -2.46958225
## [10,] -4.29884118 -4.29884118
## [11,]  6.14974070  6.14974070
## [12,]  1.48294980  1.48294980
## [13,] -2.97603592 -2.97603592
## [14,] -5.11548970 -5.11548970
## [15,] -2.86991467 -2.86991467
## [16,]  2.35779046  2.35779046
## [17,]  4.14254977  4.14254977
## [18,]  2.90230180  2.90230180
## [19,]  0.31841957  0.31841957
## [20,] -3.69962010 -3.69962009
## [21,] -0.04912758 -0.04912758
## [22,] -3.34423305 -3.34423305
## [23,] -3.36219218 -3.36219218
## [24,] -4.72380997 -4.72380997
## [25,] -0.03017356 -0.03017356
## [26,]  2.73996071  2.73996071
## [27,] -1.43443899 -1.43443899
## The above two columns you get should be very similar.
##      (Left-Your Numerical Gradient, Right-Analytical Gradient)
##
```

```
## If your implementation is correct, then
##     the relative difference will be small (less than 1e-9).
##
##     Relative Difference: 2.45535e-12
```

Before training the collaborative filtering model, we will first add ratings that correspond to a new user that we just observed. This part of the code will allow you to put in your own ratings for the movies in our dataset!

```
# Initialize my ratings
my_ratings <- rep(0,1682)

# Check the file movie_ids.txt for id of each movie in our dataset
# For example, Toy Story (1995) has ID 1, so to rate it "4", you can set
my_ratings[1] <- 4

# Or suppose did not enjoy Silence of the Lambs (1991), you can set
my_ratings[98] <- 2

# We have selected a few movies we liked / did not like and the ratings we
# gave are as follows:
my_ratings[7] <- 3
my_ratings[12] <- 5
my_ratings[54] <- 4
my_ratings[64] <- 5
my_ratings[66] <- 3
my_ratings[69] <- 5
my_ratings[183] <- 4
my_ratings[226] <- 5
my_ratings[355] <- 5

cat(sprintf('\n\nNew user ratings:\n'))
```

```
##
##
## New user ratings:
for (i in 1:length(my_ratings))
  if (my_ratings[i] > 0 )
    cat(sprintf('Rated %d for %s\n', my_ratings[i], movieList[i]))
```

```
## Rated 4 for Toy Story (1995)
## Rated 3 for Twelve Monkeys (1995)
## Rated 5 for Usual Suspects, The (1995)
## Rated 4 for Outbreak (1995)
## Rated 5 for Shawshank Redemption, The (1994)
## Rated 3 for While You Were Sleeping (1995)
## Rated 5 for Forrest Gump (1994)
## Rated 2 for Silence of the Lambs, The (1991)
## Rated 4 for Alien (1979)
## Rated 5 for Die Hard 2 (1990)
## Rated 5 for Sphere (1998)
```

Now we will train the collaborative filtering model.

```
load("movies_ratings.Rda")
list2env(data,.GlobalEnv) # "Y" and "R"
```

```

## <environment: R_GlobalEnv>
rm(data)

# Add our own ratings to the data matrix

Y <- cbind(my_ratings, Y)
R <- cbind((my_ratings != 0), R)

# Normalize Ratings
source("normalizeRatings.R")
NR <- normalizeRatings(Y, R)
Ynorm <- NR$Ynorm
Ymean <- NR$Ymean
# Useful Values
num_users <- dim(Y)[2]
num_movies <- dim(Y)[1]
num_features <- 10

# Set Initial Parameters (Theta, X)
n <- num_movies * num_features
X <- matrix(rnorm(n), num_movies, num_features)

n <- num_users * num_features
Theta <- matrix(rnorm(n), num_users, num_features)

initial_parameters <- c(c(X), c(Theta))

# Set Regularization
lambda <- 10

cF <- cofiCostFunc(Ynorm, R, num_users, num_movies, num_features, lambda)
gF <- cofiGradFunc(Ynorm, R, num_users, num_movies, num_features, lambda)

#install.packages("lbfgsb3")
library(lbfgsb3)
source("lbfgsb3_.R")
theta <- lbfgsb3_(initial_parameters, fn= cF, gr=gF, control = list(trace=1,maxit=100))$prm

## This problem is unconstrained.
## At iteration 0 f = 691521.2
## At iteration 2 f = 670660.8
## At iteration 3 f = 593480.1
## At iteration 4 f = 238705.1
## At iteration 5 f = 157828.1
## At iteration 6 f = 110506.8
## At iteration 7 f = 86787.43
## At iteration 8 f = 62428.85
## At iteration 9 f = 56515.35
## At iteration 10 f = 50350.14
## At iteration 11 f = 45894.19
## At iteration 12 f = 44464.24
## At iteration 13 f = 42849.04
## At iteration 14 f = 41895.81
## At iteration 15 f = 41350.4

```

```
## At iteration 16 f = 40771.77
## At iteration 17 f = 40396.03
## At iteration 18 f = 40085.82
## At iteration 19 f = 39881.12
## At iteration 20 f = 39742.31
## At iteration 21 f = 39584.01
## At iteration 22 f = 39528.39
## At iteration 23 f = 39436.02
## At iteration 24 f = 39386.29
## At iteration 25 f = 39326.8
## At iteration 26 f = 39273.87
## At iteration 27 f = 39232.24
## At iteration 28 f = 39205.37
## At iteration 29 f = 39178.56
## At iteration 30 f = 39152.69
## At iteration 31 f = 39133.48
## At iteration 32 f = 39119.98
## At iteration 33 f = 39096.5
## At iteration 34 f = 39085.85
## At iteration 35 f = 39072.25
## At iteration 36 f = 39059.13
## At iteration 37 f = 39048.92
## At iteration 38 f = 39043.53
## At iteration 39 f = 39028.1
## At iteration 40 f = 39022.55
## At iteration 41 f = 39012.83
## At iteration 42 f = 39009.09
## At iteration 43 f = 39001.61
## At iteration 44 f = 38998.25
## At iteration 45 f = 38994.79
## At iteration 46 f = 38991.11
## At iteration 47 f = 38987.68
## At iteration 48 f = 38986.26
## At iteration 49 f = 38984.48
## At iteration 50 f = 38982.63
## At iteration 51 f = 38981.36
## At iteration 52 f = 38980.68
## At iteration 53 f = 38979.86
## At iteration 54 f = 38979.16
## At iteration 55 f = 38978.45
## At iteration 56 f = 38978
## At iteration 57 f = 38977.54
## At iteration 58 f = 38977.21
## At iteration 59 f = 38976.86
## At iteration 60 f = 38976.45
## At iteration 61 f = 38976.24
## At iteration 62 f = 38976.02
## At iteration 63 f = 38975.76
## At iteration 64 f = 38975.52
## At iteration 65 f = 38975.32
## At iteration 66 f = 38974.91
## At iteration 67 f = 38974.72
## At iteration 68 f = 38974.41
## At iteration 69 f = 38974.01
```



```
## At iteration 70 f = 38973.57
## At iteration 71 f = 38973.2
## At iteration 72 f = 38972.71
## At iteration 73 f = 38972.46
## At iteration 74 f = 38972.15
## At iteration 75 f = 38971.74
## At iteration 76 f = 38971.46
## At iteration 77 f = 38971.15
## At iteration 78 f = 38970.57
## At iteration 79 f = 38970.31
## At iteration 80 f = 38969.93
## At iteration 81 f = 38969.85
## At iteration 82 f = 38969.28
## At iteration 83 f = 38969.08
## At iteration 84 f = 38968.73
## At iteration 85 f = 38968.22
## At iteration 86 f = 38968.47
## At iteration 87 f = 38967.89
## At iteration 88 f = 38967.39
## At iteration 89 f = 38966.98
## At iteration 90 f = 38966.54
## At iteration 91 f = 38966.16
## At iteration 92 f = 38965.68
## At iteration 93 f = 38965.17
## At iteration 94 f = 38964.75
## At iteration 95 f = 38964.36
## At iteration 96 f = 38963.8
## At iteration 97 f = 38963.26
## At iteration 98 f = 38962.78
## At iteration 99 f = 38962.13
## At iteration 100 f = 38961.38
## At iteration 101 f = 38960.87
```

```
# The following code works but optim is slow on this problem
# theta <- optim(initial_parameters, fn = cF, gr = gF,
  #method = "BFGS", control = list(maxit=10, trace=1, REPORT=1) )$par

# Unfold the returned theta back into U and W
X <- matrix(theta[1:(num_movies*num_features)], num_movies, num_features)
Theta <- matrix(theta[(num_movies*num_features+1):length(theta)], num_users, num_features)

cat(sprintf('Recommender system learning completed.\n'))
```

## Recommender system learning completed.

(2 marks) After training the model, you have now computed X and Theta. Use them to compute the top 10 recommendations for the new user. Print the movie titles and the associated ratings.

```
# ----- YOUR CODE HERE -----
p <- X %*% t(Theta)
my_predictions <- p[,1] + Ymean

movieList <- loadMovieList()

ix <- sort(my_predictions, decreasing = TRUE, index.return=TRUE)$ix
```

```

cat(sprintf('\nTop recommendations for you:\n'))

##
## Top recommendations for you:
for (i in 1:10){
  j <- ix[i]
  cat(sprintf('Predicting rating %.1f for movie %s\n', my_predictions[j],movieList[j]))
}

## Predicting rating 5.0 for movie Great Day in Harlem, A (1994)
## Predicting rating 5.0 for movie They Made Me a Criminal (1939)
## Predicting rating 5.0 for movie Santa with Muscles (1996)
## Predicting rating 5.0 for movie Saint of Fort Washington, The (1993)
## Predicting rating 5.0 for movie Prefontaine (1997)
## Predicting rating 5.0 for movie Star Kid (1997)
## Predicting rating 5.0 for movie Entertaining Angels: The Dorothy Day Story (1996)
## Predicting rating 5.0 for movie Aiqing wansui (1994)
## Predicting rating 5.0 for movie Marlene Dietrich: Shadow and Light (1996)
## Predicting rating 5.0 for movie Someone Else's America (1995)
cat(sprintf('\n\nOriginal ratings provided:\n'))

##
##
## Original ratings provided:
for (i in 1:length(my_ratings))
  if (my_ratings[i] > 0 )
    cat(sprintf('Rated %d for %s\n', my_ratings[i],movieList[i]))

## Rated 4 for Toy Story (1995)
## Rated 3 for Twelve Monkeys (1995)
## Rated 5 for Usual Suspects, The (1995)
## Rated 4 for Outbreak (1995)
## Rated 5 for Shawshank Redemption, The (1994)
## Rated 3 for While You Were Sleeping (1995)
## Rated 5 for Forrest Gump (1994)
## Rated 2 for Silence of the Lambs, The (1991)
## Rated 4 for Alien (1979)
## Rated 5 for Die Hard 2 (1990)
## Rated 5 for Sphere (1998)
# -----

```

## TURN IN

- Your .Rmd file (which should knit without errors and without assuming any packages have been pre-loaded)
- Your Word (or pdf) file that results from knitting the Rmd.
- DUE: October 2, 11:55pm (late submissions not allowed), loaded into moodle.