



ETC3555

Statistical Machine Learning

(Stochastic) gradient descent

14 August 2018

Outline

1 Gradient descent

2 Stochastic gradient descent

How to minimize E_{in}

How to minimize E_{in}

For logistic regression,

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \ln \left(1 + e^{-y_n \mathbf{w}^T \mathbf{x}_n} \right) \quad \leftarrow \text{iterative solution}$$

Compare to linear regression:

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - y_n)^2 \quad \leftarrow \text{closed-form solution}$$

Gradient descent


Iterative method: gradient descent

General method for nonlinear optimization

Start at $\mathbf{w}(0)$; take a step along steepest slope

Fixed step size: $\mathbf{w}(1) = \mathbf{w}(0) + \eta \hat{\mathbf{v}}$

What is the direction $\hat{\mathbf{v}}$?



$\hat{\mathbf{v}}$ is a unit vector, i.e. $\|\hat{\mathbf{v}}\| = 1$ and $\eta > 0$

What is the best direction $\hat{\mathbf{v}}$?

We want to minimize

$$E_{\text{in}}(\mathbf{w}(1)) - E_{\text{in}}(\mathbf{w}(0)) = E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)).$$

The Taylor expansion of $E_{\text{in}}(\mathbf{w})$ at $\mathbf{w} = \mathbf{w}(0) + \eta \hat{\mathbf{v}}$ is given by

$$E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) = E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2),$$

where $\nabla E_{\text{in}}(\mathbf{w}) = (\frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_1}, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_d})^T$ is the gradient.

We can write

$$\begin{aligned} & E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)) \\ &= \underbrace{E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2)}_{\text{Taylor approximation}} - E_{\text{in}}(\mathbf{w}(0)) \\ &\approx \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}}. \end{aligned}$$

What is the best direction $\hat{\mathbf{v}}$?

We want to minimize

$$E_{\text{in}}(\mathbf{w}(1)) - E_{\text{in}}(\mathbf{w}(0)) = E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)).$$

The Taylor expansion of $E_{\text{in}}(\mathbf{w})$ at $\mathbf{w} = \mathbf{w}(0) + \eta \hat{\mathbf{v}}$ is given by

$$E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) = E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2),$$

where $\nabla E_{\text{in}}(\mathbf{w}) = (\frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_1}, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_d})^T$ is the gradient.

We can write

$$\begin{aligned} & E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)) \\ &= \underbrace{E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2)}_{\text{Taylor approximation}} - E_{\text{in}}(\mathbf{w}(0)) \\ &\approx \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}}. \end{aligned}$$

What is the best direction $\hat{\mathbf{v}}$?

We want to minimize

$$E_{\text{in}}(\mathbf{w}(1)) - E_{\text{in}}(\mathbf{w}(0)) = E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)).$$

The Taylor expansion of $E_{\text{in}}(\mathbf{w})$ at $\mathbf{w} = \mathbf{w}(0) + \eta \hat{\mathbf{v}}$ is given by

$$E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) = E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2),$$

where $\nabla E_{\text{in}}(\mathbf{w}) = (\frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_1}, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_d})^T$ is the gradient.

We can write

$$\begin{aligned} & E_{\text{in}}(\mathbf{w}(0) + \eta \hat{\mathbf{v}}) - E_{\text{in}}(\mathbf{w}(0)) \\ &= \underbrace{E_{\text{in}}(\mathbf{w}(0)) + \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} + O(\eta^2)}_{\text{Taylor approximation}} - E_{\text{in}}(\mathbf{w}(0)) \\ &\approx \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}}. \end{aligned}$$

What is the best direction $\hat{\mathbf{v}}$?

$$\underset{\hat{\mathbf{v}}, \|\hat{\mathbf{v}}\|=1}{\text{minimize}} \quad \eta \nabla E_{\text{in}}(\mathbf{w}(0))^T \hat{\mathbf{v}} = \|\nabla E_{\text{in}}(\mathbf{w}(0))\| \|\hat{\mathbf{v}}\| \cos(\theta)$$

where θ is the angle between $\nabla E_{\text{in}}(\mathbf{w}(0))$ and $\hat{\mathbf{v}}$.

$$\equiv \underset{\hat{\mathbf{v}}}{\text{minimize}} \quad \|\nabla E_{\text{in}}(\mathbf{w}(0))\| \cos(\theta) \equiv \underset{\hat{\mathbf{v}}}{\text{minimize}} \quad \cos(\theta)$$

This quantity is minimized with $\theta = 180^\circ$ ($\cos(\theta) = -1$), i.e. $\hat{\mathbf{v}}$ is pointing in the opposite direction of the gradient. Since $\hat{\mathbf{v}}$ is a unit vector, we have $\hat{\mathbf{v}} = -\frac{\nabla E_{\text{in}}(\mathbf{w}(0))}{\|\nabla E_{\text{in}}(\mathbf{w}(0))\|}$.


In other words, we update the weights as follows

$$\mathbf{w}(1) = \mathbf{w}(0) - \eta \frac{\nabla E_{\text{in}}(\mathbf{w}(0))}{\|\nabla E_{\text{in}}(\mathbf{w}(0))\|}.$$


Which step size?

Fixed-size step?


How η affects the algorithm:



η too small



η too large



variable η – just right

η should increase with the slope

©  Creator: Yasser Abu-Mostafa - LFD Lecture 9

21/24

η should be proportional to the length of the gradient

From step size to learning rate

Easy implementation

Instead of

$$\begin{aligned}\Delta \mathbf{w} &= \eta \hat{\mathbf{v}} \\ &= -\eta \frac{\nabla E_{\text{in}}(\mathbf{w}(0))}{\|\nabla E_{\text{in}}(\mathbf{w}(0))\|}\end{aligned}$$

Have

$$\Delta \mathbf{w} = -\eta \nabla E_{\text{in}}(\mathbf{w}(0))$$

Fixed learning rate η

Gradient descent algorithm

Fixed learning rate gradient descent:

- 1: Initialize the weights at time step $t = 0$ to $\mathbf{w}(0)$.
- 2: **for** $t = 0, 1, 2, \dots$ **do**
- 3: Compute the gradient $\mathbf{g}_t = \nabla E_{\text{in}}(\mathbf{w}(t))$.
- 4: Set the direction to move, $\mathbf{v}_t = -\mathbf{g}_t$.
- 5: Update the weights: $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \mathbf{v}_t$.
- 6: Iterate to the next step until it is time to stop.
- 7: Return the final weights.

Exercise

Exercise 3.7

For logistic regression, show that

$$\begin{aligned}\nabla E_{\text{in}}(\mathbf{w}) &= -\frac{1}{N} \sum_{n=1}^N \frac{y_n \mathbf{x}_n}{1 + e^{y_n \mathbf{w}^T \mathbf{x}_n}} \\ &= \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n \theta(-y_n \mathbf{w}^T \mathbf{x}_n).\end{aligned}$$

Argue that a 'misclassified' example contributes more to the gradient than a correctly classified one.

Logistic regression algorithm

Logistic regression algorithm

- 1: Initialize the weights at $t = 0$ to $\mathbf{w}(0)$
- 2: **for** $t = 0, 1, 2, \dots$ **do**
- 3: Compute the gradient

$$\nabla E_{\text{in}} = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \mathbf{x}_n}{1 + e^{y_n \mathbf{w}^T(t) \mathbf{x}_n}}$$

- 4: Update the weights: $\mathbf{w}(t+1) = \mathbf{w}(t) - \eta \nabla E_{\text{in}}$
- 5: Iterate to the next step until it is time to stop
- 6: Return the final weights \mathbf{w}

Stopping criterion


Typically the initial point $\mathbf{w}(0)$ is picked randomly, or we use prior knowledge about the problem. But when to stop the algorithm?

Some common choices (ϵ is a small prescribed threshold):

- $\|\nabla E_{\text{in}}(\mathbf{w}(t))\| < \epsilon$
- $|E_{\text{in}}(\mathbf{w}(t+1)) - E_{\text{in}}(\mathbf{w}(t))| < \epsilon$
- $\|\mathbf{w}(t+1) - \mathbf{w}(t)\| < \epsilon$
- $\frac{|E_{\text{in}}(\mathbf{w}(t+1)) - E_{\text{in}}(\mathbf{w}(t))|}{\max\{1, |E_{\text{in}}(\mathbf{w}(t))|\}} < \epsilon$
- $t > T$

Choosing the learning rate

- The size of the step taken in gradient descent, $-\eta \nabla E_{\text{in}}$, is proportional to the learning rate η .
- The optimal step size/learning rate depends on how wide or narrow the error surface is near the minimum.
- Wider surface \implies we can take larger steps without overshooting. Since $\|\nabla E_{\text{in}}\|$ is small, we need a large η .



Choosing the learning rate

Variable Learning Rate Gradient Descent:

- 1: Initialize $\mathbf{w}(0)$, and η_0 at $t = 0$. Set $\alpha > 1$ and $\beta < 1$.
- 2: **while** stopping criterion has not been met **do**
- 3: Let $\mathbf{g}(t) = \nabla E_{\text{in}}(\mathbf{w}(t))$, and set $\mathbf{v}(t) = -\mathbf{g}(t)$.
- 4: **if** $E_{\text{in}}(\mathbf{w}(t) + \eta_t \mathbf{v}(t)) < E_{\text{in}}(\mathbf{w}(t))$ **then**
- 5: accept: $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta_t \mathbf{v}(t)$; $\eta_{t+1} = \alpha \eta_t$.
- 6: **else**
- 7: reject: $\mathbf{w}(t+1) = \mathbf{w}(t)$; $\eta_{t+1} = \beta \eta_t$.
- 8: Iterate to the next step, $t \leftarrow t + 1$.

Steepest Descent (Gradient Descent + Line Search):

- 1: Initialize $\mathbf{w}(0)$ and set $t = 0$;
- 2: **while** stopping criterion has not been met **do**
- 3: Let $\mathbf{g}(t) = \nabla E_{\text{in}}(\mathbf{w}(t))$, and set $\mathbf{v}(t) = -\mathbf{g}(t)$.
- 4: Let $\eta^* = \operatorname{argmin}_{\eta} E_{\text{in}}(\mathbf{w}(t) + \eta \mathbf{v}(t))$.
- 5: $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta^* \mathbf{v}(t)$.
- 6: Iterate to the next step, $t \leftarrow t + 1$.

Other optimization methods

- Momentum, Nesterov Momentum, ...
- Adaptive learning rates: AdaGrad, RMSProp, RMS Prop, Adam
- Newton's Method, Conjugate gradient, BFGS, L-BFGS

Outline

1 Gradient descent

2 Stochastic gradient descent

About gradient descent

- Computing the full gradient is slow for big data
- Stuck at stationary points

Stochastic gradient descent

Stochastic gradient descent

GD minimizes:

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \underbrace{e(h(\mathbf{x}_n), y_n)}_{\ln(1 + e^{-y_n \mathbf{w}^T \mathbf{x}_n})} \leftarrow \text{in logistic regression}$$

by iterative steps along $-\nabla E_{\text{in}}$:

$$\Delta \mathbf{w} = -\eta \nabla E_{\text{in}}(\mathbf{w})$$

∇E_{in} is based on all examples (\mathbf{x}_n, y_n)

“batch” GD

The stochastic aspect

The stochastic aspect

Pick one (\mathbf{x}_n, y_n) at a time. Apply GD to $\mathbf{e}(h(\mathbf{x}_n), y_n)$

“Average” direction:

$$\mathbb{E}_n [-\nabla \mathbf{e}(h(\mathbf{x}_n), y_n)] = \frac{1}{N} \sum_{n=1}^N -\nabla \mathbf{e}(h(\mathbf{x}_n), y_n)$$
$$= -\nabla E_{\text{in}}$$

randomized version of GD

stochastic gradient descent (SGD)



Benefits of SGD

Benefits of SGD

1. cheaper computation
2. randomization
3. simple

Rule of thumb:

$\eta = 0.1$ works



randomization helps

Exercise 3.10

- (a) Define an error for a single data point (\mathbf{x}_n, y_n) to be

$$e_n(\mathbf{w}) = \max(0, -y_n \mathbf{w}^T \mathbf{x}_n).$$

Argue that PLA can be viewed as SGD using e_n .

- (b) For logistic regression with a very large \mathbf{w} , argue that minimizing E_{in} using SGD is similar to PLA. This is another indication that the logistic regression weights can be used as a good approximation for classification.

Mini-batch gradient descent

Compute the gradient using $1 \leq b \leq N$ samples.

- 1 Pick b examples ($1 \leq b \leq N$)
- 2 Apply batch GD to these b examples

Note: we can also shuffle the data and pick the mini-batches sequentially.

- $b = N$ is GD and $b = 1$ is SGD
- Bias and variance tradeoff
- A single pass through the entire training data is called an *epoch*. With mini-batches of size b , we update the parameters N/b times per epoch.
- We often need multiple epochs to obtain a good training accuracy.